
Integrated Longitudinal Data Visualization Platform for INDEPTH HDSS Sites



Tumaini Kilimba

A Dissertation submitted to the Faculty Of Health Sciences for the partial
fulfilment of the requirements for the degree of

Master of Science

Supervisors:

Gideon Nimako¹, Kobus Herbts²

¹School of Public Health
Faculty of Health Sciences
University of the Witwatersrand

²Africa Centre for Health and Population Studies
Mtubatuba
Kwa-Zulu Natal

Copyright © by
Tumaini Kilimba
2015

DECLARATION

I declare that this dissertation is my own, unaided work under the supervision of Gideon Nimako and Kobus Herbts. It is being submitted for the Degree of Masters of Science at the University of the Witwatersrand, Johannesburg. It has not been submitted before for any degree or examination at any other University.

22nd July, 2015

ABSTRACT

A lot of work has been done on spatial visualizations for public health. The feasibility of creating poverty maps for Indonesia at various administrative levels to help with the implementation of programs which target the poor was investigated in [6]. Their focus however was on a new methodology for imputing per capita consumption for each household in the population based on data collected from household surveys and data collected from population censuses. Their report however did not focus much on how the visualization platform was to be built. The final product though, after all the computations is the visualization, the poverty map of the country.

ACKNOWLEDGEMENTS

XXX

CONTENTS

1	Introduction	1
1.1	Problem Statement	3
1.2	Motivation	3
2	Background and Related Works	4
3	Research Aims and Objectives	8
3.1	Problem Statement	9
3.2	Motivation	9
3.3	Overall Aim	10
3.4	Specific Objectives	10
4	HDSS Sites Visualization Platform	11
4.1	Scope of Research	12
4.2	Overview of the Solution	12
4.2.1	Scientific Portal	12
4.2.2	Data Operations Portal	13
4.2.3	Community Engagement Portal	14
5	Experimental Setup, Models and Validation	16
5.1	Scientific Portal	17
5.1.1	Outcome Pyramid	17
5.1.2	Time Series	18
5.1.3	Image Plot (Heat Map	18
5.2	Implementation of the Science Portal	21
5.2.1	Data Preparation	21

6

Chapter Title

24

6.1

Sub Title 1

25

6.2

Sub Title 2

25

7

Conclusion and Future Directions

26

7.1

Summary of Research Contributions

27

A

Appendix

28

A.1

XXX

28

References

29

Dedicated to

XXX

AAA and

BBB

INTRODUCTION

The first chapter introduces the subject of visual analytics in general and then goes on to introduce visual analytics in the context of public health and epidemiology. We will delve into the literature around this subject and then we will espouse our own motivations in creating a generalizable public health data visual platform

Contents

1.1	Problem Statement	3
1.2	Motivation	3

Visual analytics tools are used to synthesize information and glean insight from massive, dynamic, ambiguous and often conflicting data. Visual Analytics (VA) is often referred to as a means for dealing with complex, large information sources that require human judgment to identify the expected and discover the unexpected. It is a multidisciplinary field whose core areas are analytical reasoning techniques, visual representation and interaction techniques, data representations and transformations as well as production, presentation and dissemination [1].

Analytic reasoning techniques enable users to get deep insights which support assessment, planning and decision making. It also allows users to assimilate large amounts of information at once [2]. In a healthcare setting expected outcomes include more efficient and effective clinical performance monitoring and improvement as well as improved modelling of patient flow and management. Additionally, one can expect increased quality of care, improved safety and efficiency and better support for clinical costing and resource coordination. It can also lead to better planned growth and competitive advantage [cite 3].

VA is referred to as “the science of analytical reasoning facilitated by interactive visual interfaces”. Map based community health visualizations have provided a comprehensive and powerful interface for scientists and policy makers to visualize health care quality, public health outcomes and access to care. This has helped in making evidence-based decisions about improving healthcare [cite 4-7]

Multi-panel graphs have also been used to good effect in a graphical tool for epidemiological studies to reveal the distribution of an outcome by time and age simultaneously [cite 8].

The growth of surveillance systems in both quantity of data and variety of outcomes is likely to necessitate constant innovations in data processing, synthesis, and communication [cite 8] Therefore, techniques to support production, presentation and dissemination of analytic results will allow us to communicate the information in the appropriate context to a variety of audiences.

One critical requirement for successful public health surveillance is the ability to analyse and present data so that it is understandable to a range of public health stakeholders. In public health, VA can be viewed as the bridge between the availability of surveillance data in database architectures and useful information derived from this available data [cite 9].

INDEPTH Health and Demographic Surveillance Sites (HDSS's) deal with complex longitudinal data and, as a result, knowledge transfer to stakeholders is challenging. Better visualisation of this data is therefore required in order for potential scientific users to maximise exploratory data analysis and hypothesis generation. It would also aid decision-makers and the society at large to visualise this information in terms understandable by them. Such a visualization tool will also improve field work research activities by providing summary data of operational progress, e.g. fieldwork data collection progress, data entry progress, or other parameters such as data quality. This will serve to improve operational decision making and data quality. However, datasets at HDSS sites are normally under-visualised. These HDSS sites currently have no generalizable framework for implementing a data visualization platform to be used at these sites.

The current under-visualization of HDSS datasets shows little promise of improvement in a harmonised way (across multiple sites) unless specific research efforts are directed towards finding a generalizable solution for delivering interactive visualizations, supporting exploratory analyses and real-time displays of operational progress. Furthermore, there is a paucity of research specifically on the technologies and tools which can be used to create such a data visualization [cite 10].

BACKGROUND AND RELATED WORKS

This chapter gives the background and related works that have been done in the field of public health data visualization

The topic of visualization of public health data was identified in 2009 by the CDC as one of six major concerns which must be addressed by the public health community in order to advance public health surveillance in the 21st century [cite 9]. However, very little has been researched in terms of standardization of the workflow and linking technologies for heterogeneous data sources needed specifically for visualizations in public health science [cite 10].

Of crucial importance when dealing with large datasets is the need for the users of the data, be they scientists or any other stakeholders to be able to discover the relations among and between the results of data analyses and queries [cite 10]. However, due to bottlenecks resulting from resource cost and lack of required skills, data visualization becomes an end product of scientific analysis rather than an exploration tool which facilitates scientists to generate better hypotheses in the continually more data-intensive scientific process [cite 10]. The use of such visualizations are usually utilised in business analytics, open government data systems, and media infographics but have generally not been used in public health. However, the capabilities currently being seen by web-based tools and technologies may be the breakthrough in resolving the scientific visualization bottleneck.

A lot of work has been done on spatial visualizations for public health [cite 4,6,7]. The feasibility of creating poverty maps for Indonesia at various administrative levels to help with the implementation of programs which target the poor was investigated in [cite 6]. Their focus however was on a new methodology for imputing per capita consumption for each household in the population based on data collected from household surveys and data collected from population censuses. Additionally, their report did not focus much on how the visualization platform was to be built. The final product though, after all the computations is the visualization, the poverty map of the country.

Along a similar spatial theme, [cite 7] looked at how the agents of parasitic diseases are spatially distributed using map visualizations. The tools of interest used in their research were the two closely linked Google products, Google Earth and Google Maps. Though they provide a little

more implementation details, the mix of tools used are not all open-source.

Other research on implementing a data visualization platform for community health assessment (CHA) used open source technologies but was limited to a desktop application and could not be accessed online [cite 11]. Web based tools have been seen as the preferred platform of choice for public health researchers as they permit distributed access, reduced software implementation costs and wider exposure of public health information for public dissemination [cite 4,12,13]. The report also gives scant details of the actual open source tools used and how they were put together in a way which would allow recreation of the steps.

The use of multi-panel graphs to illustrate trends and anomalies which would otherwise be obscured by traditional epidemiological visualization techniques such as pyramids and time-series plots was explored in [cite 8]. Under future developments in their report however, they acknowledge that two other features if incorporated to these graphs would enhance their impact, namely the dynamic display of data and interactivity.

Existing tools offer a range of features and functions to allow for exploration, analysis and visualization of public health users data, but the tools are often for siloed applications incapable of reciprocal operation with other, related information systems. They are isolated to the jurisdictions and organisations which developed them limiting their widespread adoption by other agencies or organisations[cite 13]. Interoperability of the visualization tools has been identified in a systematic literature review of infectious disease visualization tools as a prominent theme, due to challenges associated with increasingly collaborative and interdisciplinary nature of disease surveillance, control and prevention [cite 13].

The CDC's inclusion of VA as one of its six areas of focused research in public health lends credence to the importance of this research project. Furthermore there is little that has been researched in terms of standardisation of the workflow, tools and linking technologies for visualizations in the public health domain, as well as few existing web based tools for health related data visualization [cite 4,10]. This research will add to the body of literature broadly

in the subject of tools and technologies for data visualization in general, and specifically for implementing a data visualization platform which is generalizable for all INDEPTH HDSS sites with the aim of allowing it to become a central piece of the scientific process. It will also augment on work already done [cite 8] by incorporating dynamic display of data and interactivity into multi-panel graphs for epidemiological research [cite 8].

RESEARCH AIMS AND OBJECTIVES

In this chapter we will elaborate on the problem we are addressing. We shall also discuss how this problem is relevant to Health and Demographic Surveillance Sites (HDSS's) and our motivation to address it

Contents

3.1	Problem Statement	9
3.2	Motivation	9
3.3	Overall Aim	10
3.4	Specific Objectives	10

3.1 Problem Statement

One critical requirement for successful public health surveillance is the ability to analyse and present data so that it is understandable to a range of public health stakeholders. In public health, VA can be viewed as the bridge between the availability of surveillance data in database architectures and useful information derived from this available data [cite 9].

INDEPTH Health and Demographic Surveillance Sites (HDSS's) deal with complex longitudinal data and, as a result, knowledge transfer to stakeholders is challenging. Better visualisation of this data is therefore required in order for potential scientific users to maximise exploratory data analysis and hypothesis generation. It would also aid decision-makers and the society at large to visualise this information in terms understandable by them. Such a visualization tool will also improve field work research activities by providing summary data of operational progress, e.g. fieldwork data collection progress, data entry progress, data archiving progress or other parameters such as data quality. This will serve to improve operational decision making and data quality. However, datasets at HDSS sites are normally under-visualised. These HDSS sites currently have no generalizable framework for implementing a data visualization platform to be used at these sites.

3.2 Motivation

The current under-visualization of HDSS datasets shows little promise of improvement in a harmonised way (across multiple sites) unless specific research efforts are directed towards finding a generalizable solution for delivering interactive visualizations, supporting exploratory analyses and real-time displays of operational progress. Furthermore, there is a paucity of research specifically on the technologies and tools which can be used to create such a data visualization [cite 10].

3.3 Overall Aim

The overall aim of the project was to increase the utilization of data in INDEPTH sites through interactive data visualization. This was aimed at improving hypotheses generation at these research sites as well as increase operational awareness.

3.4 Specific Objectives

The specific objectives of this research were:

1. To design a data visualization platform for the Africa Centre for Health and Population Studies (ACHPS).
2. To build a data visualization platform for ACHPS in order to increase data utilization and hypotheses generation at the site
3. To create a developer manual for data visualization so that the process for building the platform can be reproduced.
4. iv. To evaluate the usability of the developed platform for easy integration into the operational research cycle of the site

HDSS SITES VISUALIZATION

PLATFORM

This chapter gives a description of the the scope of the research, and overview of the solution, including a discussion on the three themes of Scientific Portal, Data Operations Portal and Community Engagement Portal. It goes on to describe in detail how the the application which serve the three themes were constructed, in terms of the tools, programming platforms and languages as well as relevant algorithms to replicate key artefacts.

Contents

4.1	Scope of Research	12
4.2	Overview of the Solution	12
4.2.1	Scientific Portal	12
4.2.2	Data Operations Portal	13
4.2.3	Community Engagement Portal	14

4.1 Scope of Research

The scope of this research is limited to the integration, transformation and visualization of datasets from demographic surveillance data, verbal autopsy data [cite for VA] and socio-economic status data.

4.2 Overview of the Solution

The Africa Centre has embarked on a concept project dubbed Data Everywhere. Its aim is to increase the comprehension, access and utility of data collected through the use of a data visualization platform with 3 themes; Scientific Portal, Data Operations Portal and Community Engagement Portal. On site, this will be realised through the placement of three 52 inch touch screens in three strategic positions depending on the target audience. These touch screens will allow for users to visually interact with data on demand, selecting, filtering and visual feedback on being touched (active assimilation), as well as animate visualizations to show temporal trends when on standby mode, allowing for passive assimilation of potential insights. Additionally, the Scientific Portal and Community Engagement Portal will be hosted on a web server which allows access to their respective visualizations remotely from a browser.

4.2.1 Scientific Portal

This screen is placed strategically in the “Science Lounge” at the Africa Centre, a lounge area where Africa Centre scientists congregate for informal discussions and coffee breaks. It will allow scientists to either passively glean insights from the wall mounted 52 inch touch-screen as the animations show trends through time, or engage with the visualizations by directly interacting with the visualizations through selections, filtering and dynamic visual feedback.

The aim of this portal is to facilitate scientific discourse, insight generation and hypothesis formulation either serendipitously (passive) or through deliberate interaction with the visu-

alizations (active), amongst Africa Centre scientists using the lounge. Furthermore, as the Scientific Portal is web hosted, it allows for external scientists who potentially want to run or collaborate on studies at the Africa Centre to get a quick feel of not only what kind of data the Africa Centre currently has but also what the data is saying in a well packaged and easily accessible manner.

The solution builds on work already done by creating multi-panel graphs with dynamic display of data and interactive capabilities [cite 8]. These graphs weave together temporality and demographics and they use time-series plots, image plots and outcome pyramids.

This portal relies on Africa Centre's demographic surveillance data. The first step in developing the visualizations for this portal was to create Extract, Transform and Load (ETL) transformations using Pentaho Kettle. These create and store the dataset for each indicator by pulling data from a Microsoft SQL Server database. An additional transformation was produced for creating/updating a lookup file which links indicators to their respective datasets via a Uniform Resource Identifier (URI). These transformations are then integrated into a single Kettle Job. The motivation for this is flexibility, as any new indicators to be visualised in the future simply need a transformation for creating the appropriate dataset and a new entry in the lookup file linking the new indicator to its dataset. This ensures that no additional programming is required on the data visualization application with each new indicator to be visualised as long as the datasets stick to predefined structural and naming conventions.

4.2.2 Data Operations Portal

The Data Operations Portal screen is placed within the Data Centre. This is the office which deals with data collection, data entry, data cleaning, data quality assurance and data archiving. This will be a dashboard which aims to give real-time feedback on the progress of data collection activities, data entry activities and data archiving activities measured against the total number of data forms allocated for a particular survey round of three studies; Household, Individual

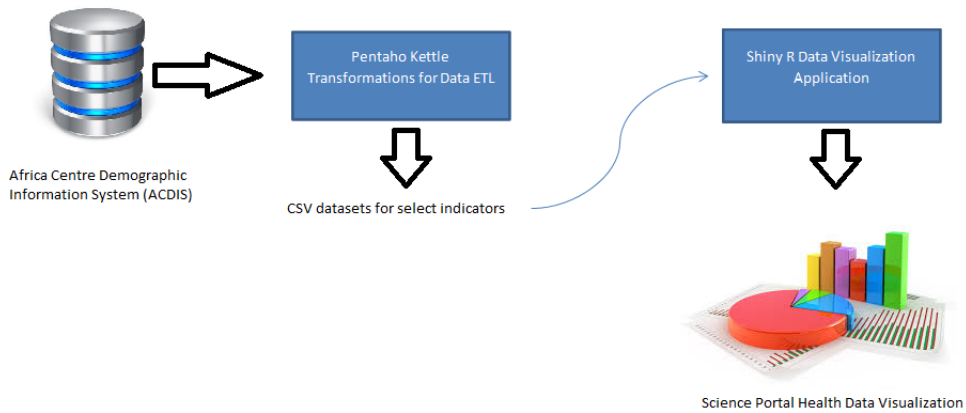


Figure 4.2.1: Overview of Scientific Portal solution

and Verbal Autopsy. This is in order to keep the Data Centre team abreast of their progress and operational bottlenecks in a transparent and accessible manner.

For this portal, data will be pulled directly from a Microsoft SQL Server via polling for changes in the underlying data every hour. This portal is not be externally accessible as it is purely for operational monitoring.




Figure 4.2.2: Overview of Data Operations Portal solution

4.2.3 Community Engagement Portal

The third and final screen is strategically placed in the Africa Centre foyer, visible and accessible to both staff and visitors. Its main focus is to package data which is of interest to the community which Africa Centre's research serves into visual representations that are easy

to interpret. This portal can be externally accessible over the internet for the community at large to access.

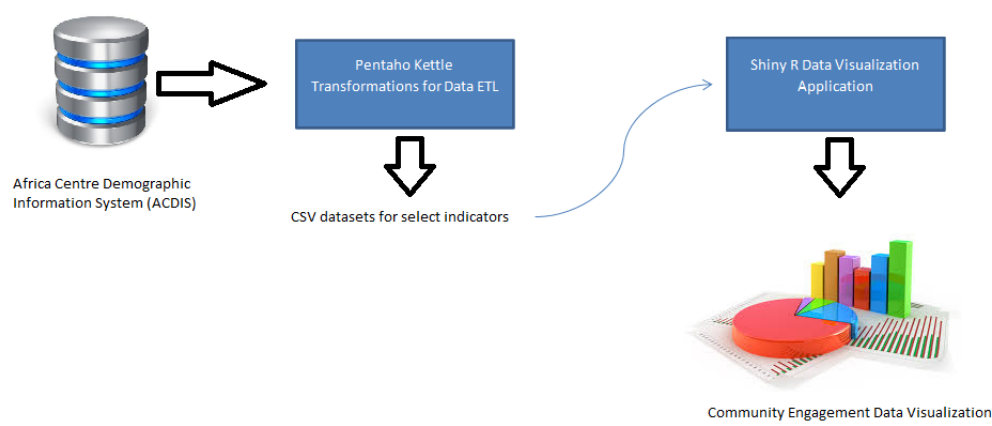


Figure 4.2.3: Overview of Community Engagement Portal solution

This research is about creating a generalizable data visualization solution for INDEPTH HDSS sites, and the output will be a software application which can visualize data at any of these sites. Different sites have differing data cleaning and data manipulation strategies and as such, the onus of ensuring that the data fed into the visualization platform is clean and has handled missing data falls on the data manager at the site. The only constant in the provided solution will be the software application; the ETL transformations and jobs at each INDEPTH site will have to be developed by the site data managers to handle each HDSS sites’ database idiosyncrasies which are best known by the on site data manager. In order to do this, they will be informed by certain dataset structures and conventions which we have documented and are elaborated on in chapter 5.

EXPERIMENTAL SETUP, MODELS AND VALIDATION

In this section we look at how we will set up our proposed solution, so that it can be replicated, as well as how we propose to validate the framework developed. We will look in detail at each of the three visualization themes (science, operations and community) in isolation.

Contents

5.1	Scientific Portal	17
5.1.1	Outcome Pyramid	17
5.1.2	Time Series	18
5.1.3	Image Plot (Heat Map)	18
5.2	Implementation of the Science Portal	21
5.2.1	Data Preparation	21

The Scientific Portal is the cornerstone of the Data Everywhere project and is based on work already done in previous research and is primarily intended to convey trends and anomalies which would otherwise be obscured by traditional visualization techniques such as case pyramids and time series plots [cite 8]. The authors of that research propose the use of multi-panel (MP) graph to achieve this. The MP graph consists of an outcome pyramid, a time-series plot and an image plot (heat map) positioned on a single canvas where at least two or more graphs share at least one common axis. This design allows for the simultaneous visualization of population structure and temporal trend.

5.1.1 Outcome Pyramid

An outcome pyramid is similar to a population pyramid, a type of graph used to describe the age and sex composition of a population. An outcome pyramid describes the age and sex distribution of a population who have a particular health or disease outcome. The pyramid itself is basically two vertically juxtaposed histograms, one for males and another for females and a vertical axis common to the two histograms which displays age. This can be represented as single years or 5 year categories. The horizontal axis can display absolute counts or rates. The longer a bar extends from the vertical axis the greater the proportion or number of individuals in that age category.

When looking at outcome pyramids, three things should be deserving of attention: the overall shape of the pyramid, such as uniform, triangular or even inverted triangle; the presence or absence of symmetry by gender; and irregularities such as bumps or spikes in the cases [cite 8]. In of themselves, outcome pyramids provide a useful tool for assessing age and gender disparities in diseases. However, as noted in [cite 8], static outcome pyramids fail to display temporal changes, which can be quite extensive especially for infectious diseases with well pronounced seasonality.

5.1.2 Time Series

Time-series plots are used to visualise temporal trends and seasonal patterns in disease rates. For the time series to be informative, it must span a sufficiently long time frame and a carefully chosen unit of time (such as day, week, month or year). For many outbreak surveillance systems, a weekly time series is standard [insert citation]. However, INDEPTH sites were not structured for outbreak surveillance, rather they excel at trend surveillance, and therefore it makes more sense to choose a year as a unit of surveillance time in our case. The vertical axis in the time series can show either disease counts or rates.

Although a well-constructed time series plot can show the temporal fluctuations in disease occurrence, interpretation of the fluctuations requires an assumption of equally distributed risk through the represented ages in the population. In reality diseases seldom affect all ages uniformly. This erroneous assumption can be partially alleviated by compiling time series plots for specific age groups, choosing a medium sized interval (5 year) [cite 8]. The added interactivity afforded by our web application allowed us to choose this option allowing a user to choose an age group of interest to visualise on the time series, thus adding a demographic dimension to it, whereas [cite 8] had a disaggregated time-series with no demographic dimension.

5.1.3 Image Plot (Heat Map)

An image plot or heat map is a data visualization tool which can be used to graphically represent your data for at least three variables. At a basic level, they can be thought of as spreadsheets which have colours instead of numbers, whereby the colour of each cell or rectangle is related to the magnitude (count or rate) of the cell amount [cite 14]. In our case the first two variables age (age groups) and time (year) are represented by the vertical axis and horizontal axis respectively. The third variable is outcome count/rate (outcome being the disease or health outcome of interest) and it is represented by the different hues or saturation of colours within the plot.

New information previously obscured by outcome pyramids and time series plots can surface on using image plots, although reading the graph properly requires some training and practice [cite 8]




Figure 5.1.1: Typical patterns observed in image plots used to study the association between age, time and the disease/outcome of interest [cite 8]

Some typical patterns are shown in 5.1.1. Panel A shows a decrease in colour from left to right indicating a decrease in the outcome of interest over time, which is more or less uniform among all ages.

Panel B shows a decrease in saturation from top to bottom showing an increase in relative risk of the outcome in older age groups irrespective of time. Another variation of B is its' inverse, whereby the saturation is darker in the lower age groups and lighter in the older age groups, signifying increased relative risk in the younger age groups for that particular outcome.

Panel C is a combination of both Panel A and Panel B in that the older age groups constantly remain at highest risk of the outcome of interest but this risk is gradually reducing over all age groups over time.

Panel D shows a striated pattern which is evidence of seasonality. The outcome of interest increases and wanes more or less evenly over all age groups periodically.

Panel E is a striated pattern which suggest a phenomenon known as the “age-cohort effect“. This is where there is an observable variation in the risk of a health outcome according to the year of birth [cite 15]. It is seen when an outcome is measured in the same cohort repeatedly and the outcome remains higher or lower in a particular group of subjects as they aged.

Lastly, panel F overall shows a similar pattern to pattern B whereby the outcome is increasingly

higher as you go up the age groups, but relatively unchanging over time. However the four distinct clusters in the younger age group seen at specific periods in time within the panel could indicate an outbreak, due to the disproportionately high values of the outcome of interest.

The combination of these three types of plots, outcome pyramid, time series and heat map strategically juxtaposed onto a single canvas give us an MP graph (5.1.2 and 5.1.3).

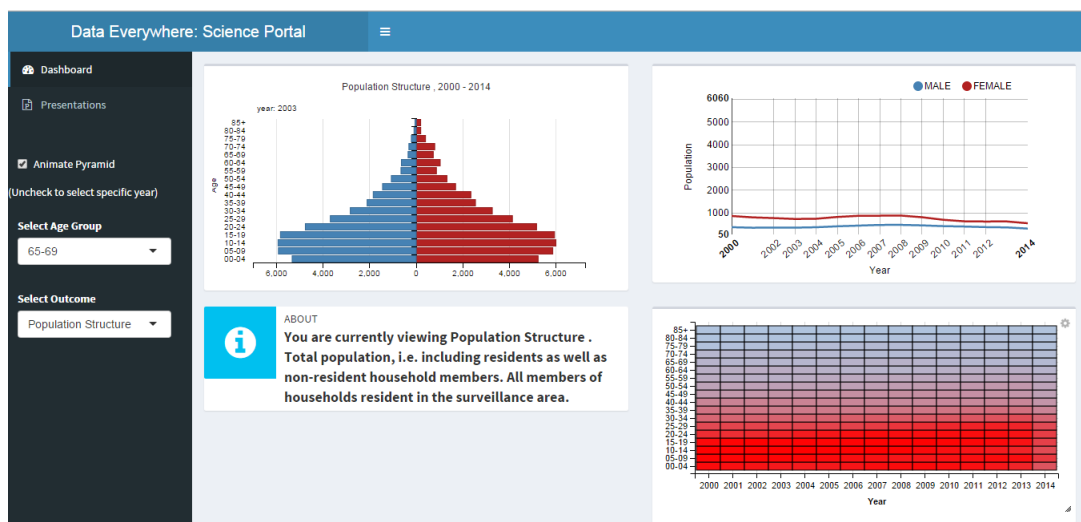


Figure 5.1.2: MP graph showing the population structure of a population in a particular health and demographic surveillance site

In 5.1.2, the top left panel is an outcome pyramid which is set to dynamically change over time to animate the changes in the age and sex distribution. The top right panel is a time series showing the trend in the population disaggregated by sex for a particular age group (here the age group is 65-69, selected from a drop down menu on the left sidebar). The lower right panel shows a heat map which indicates a higher population density in the lower age groups.

In 5.1.3, we see the same MP canvas showing a different outcome of interest, in this case All Cause Mortality Rate of the population at the Africa Centre Health and Demographic Surveillance System site. The outcome of interest is selected from a drop down menu on the left sidebar (in this graph it is visibly showing 5 possible outcomes to choose from, Population Structure, Resident Population Structure, All Cause Mortality Rates, All Cause Mortality Rates (Resident) and HIV Prevalence Rates).

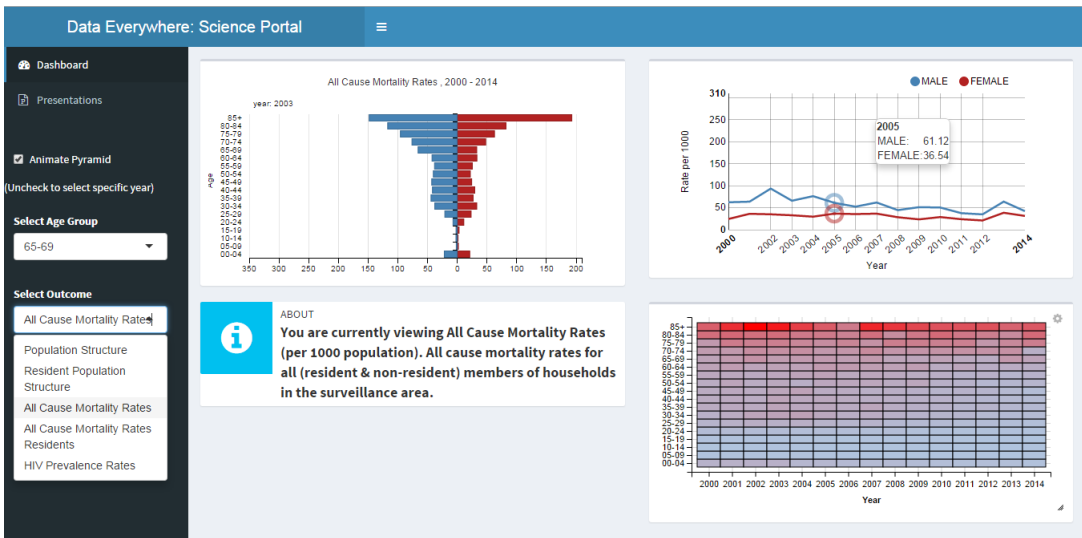


Figure 5.1.3: MP graph showing All Cause Mortality Rate of a population in a particular health and demographic surveillance site

5.2 Implementation of the Science Portal

In this section we look at how we implement the Scientific Portal, from the data preparation, the programming algorithms used and the tools and technologies involved.

5.2.1 Data Preparation

The emphasis for the scientific portal was to create a solution which was generalizable enough to be moved to other INDEPTH sites and be integrated little to no modification to the software. The key to this was to make the scientific portal metadata driven, while following the data structures and conventions which we expand on here. We use Pentaho Kettle transformations and jobs to perform all ETL tasks described here.

The data is prepared in the three steps summarized below:

1. Prepare an empty indicator lookup file
2. Add an indicator x entry into the indicator lookup file which includes in one of its fields the intended URI of the dataset for the indicator

3. Generate the data for the indicator x and place the data set at the URI indicated in step 2

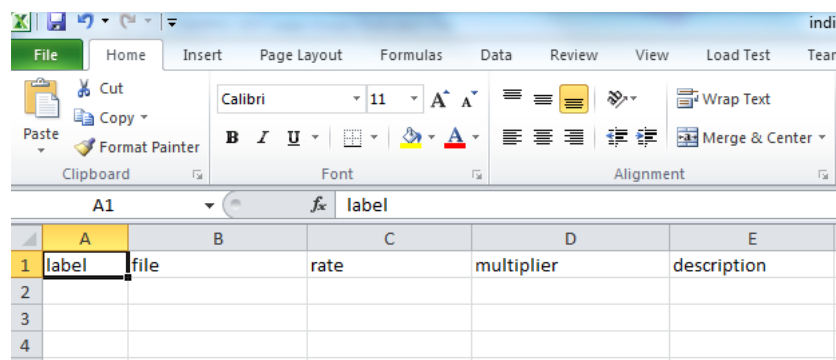
These steps are repeated for each new indicator which we would wish to visualise on the Scientific Portal. This is accomplished with a Pentaho Kettle transformation for each step, which is eventually incorporated into a Pentaho Kettle job for a single point of execution.

5.2.1.1

Prepare indicator file

The first step involves creating an empty indicator file as a Comma Separated Variable (csv) file. This is the lookup file for all indicators which we wish to visualize. The file follows a naming and structural convention which has to be followed as the application depends on these conventions. The file has the fields named *label*, *file*, *rate*, *multiplier* and *description*.

The *label* field will hold the name of the indicator and will be used to populate the drop down menu on the user interface. The *file* field will hold the URI location of the indicator and will be used to tell the application where to find the data related to a selected choice in the drop down menu. The *rate* field will hold a Boolean value flagging whether the indicator is a rate or not. The *multiplier* field is the constant value with which to multiply a calculated value if the indicator is flagged as a rate. The *description* field will hold some descriptive text describing the indicator selected in the drop down menu.



The screenshot shows the Microsoft Excel interface with a new workbook. The ribbon at the top includes File, Home, Insert, Page Layout, Formulas, Data, Review, View, Load Test, and Tear. The Home ribbon is active, showing options for Clipboard, Font, and Alignment. The font is set to Calibri, size 11. The alignment section shows options for text alignment and wrapping. The worksheet grid shows a single row of headers in row 1: 'label' in column A, 'file' in column B, 'rate' in column C, 'multiplier' in column D, and 'description' in column E. The cells are empty for rows 2, 3, and 4.

	A	B	C	D	E
1	label	file	rate	multiplier	description
2					
3					
4					

Figure 5.2.1: The initial empty indicator lookup file with the required fields

6

CHAPTER TITLE

This chapter describes xx

Contents

6.1	Sub Title 1	25
6.2	Sub Title 2	25

6.1

Sub Title 1

6.2

Sub Title 2

CONCLUSION AND FUTURE DIRECTIONS

This chapter concludes the thesis and gives roadmap for future work

Contents

7.1 Summary of Research Contributions	27
---	----

A

APPENDIX

A.1	XXX
-----	-----

REFERENCES

- [1] J. J. Thomas, K. Cook *et al.*, “A visual analytics agenda,” *Computer Graphics and Applications, IEEE*, vol. 26, no. 1, pp. 10–13, 2006.
- [2] K. A. Cook and J. J. Thomas, “Illuminating the path: The research and development agenda for visual analytics,” Pacific Northwest National Laboratory (PNNL), Richland, WA (US), Tech. Rep., 2005.
- [3] A. Sopan, A. S.-I. Noh, S. Karol, P. Rosenfeld, G. Lee, and B. Shneiderman, “Community health map: A geospatial and multivariate data visualization tool for public health datasets,” *Government Information Quarterly*, vol. 29, no. 2, pp. 223–234, 2012.